

BA assignment 3 Regression

Pratheek Sreerangam


2022-11-12

1)

```
set.seed(2017)
X=runif(100)*10
Y=X*4+3.45
Y=rnorm(100)*0.29*Y+Y
## a)
cor(X,Y)
```

```
## [1] 0.807291
```

```
plot(X,Y)
```



```
## Yes, we can fit a linear model with a positive correlation between y and x.
## b)
model<-lm(Y~X)
summary(model)
```


```
##
## Call:
## lm(formula = Y ~ X)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -26.755  -3.846  -0.387   4.318  37.503
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   4.4655     1.5537   2.874  0.00497 **
## X             3.6108     0.2666  13.542 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 7.756 on 98 degrees of freedom
## Multiple R-squared:  0.6517, Adjusted R-squared:  0.6482
## F-statistic: 183.4 on 1 and 98 DF,  p-value: < 2.2e-16
```

```
## The accuracy of the above mentioned linear model is 65.17 percent, and the variability in y may be
##  $Y=3.6108X+4.4655$  is the model's equation.
## c)
(cor(Y,X))^2
```


```
## [1] 0.6517187
```

```
## The square of a correlation has more than one r-square.
## (Correlation Coefficient)^2=Coefficient of Determination
## d)
hist(model$residuals,freq = FALSE,ylim = c(0,0.05))
```

Histogram of model\$residuals



```
qqnorm(model$residuals,col="red")
```



The linear model is acceptable since the above mentioned graph shows that residuals are typically d

2a)

```
head(mtcars)
```

```
##           mpg  cyl  disp  hp  drat    wt  qsec vs  am  gear  carb
## Mazda RX4      21.0   6  160  110 3.90 2.620 16.46 0  1    4    4
## Mazda RX4 Wag  21.0   6  160  110 3.90 2.875 17.02 0  1    4    4
## Datsun 710     22.8   4  108   93 3.85 2.320 18.61 1  1    4    1
## Hornet 4 Drive  21.4   6  258  110 3.08 3.215 19.44 1  0    3    1
## Hornet Sportabout 18.7   8  360  175 3.15 3.440 17.02 0  0    3    2
## Valiant        18.1   6  225  105 2.76 3.460 20.22 1  0    3    1
```

```
summary(lm(hp~wt,data=mtcars))
```

```
##
## Call:
## lm(formula = hp ~ wt, data = mtcars)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -83.430 -33.596 -13.587   7.913 172.030
##
```

```
## Coefficients:
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -1.821      32.325  -0.056   0.955
## wt           46.160       9.625   4.796 4.15e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 52.44 on 30 degrees of freedom
## Multiple R-squared:  0.4339, Adjusted R-squared:  0.4151
## F-statistic:    23 on 1 and 30 DF,  p-value: 4.146e-05
```

```
summary(lm(hp~mpg,data=mtcars))
```

```
##
## Call:
## lm(formula = hp ~ mpg, data = mtcars)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -59.26 -28.93 -13.45  25.65 143.36
##
## Coefficients:
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept)  324.08      27.43  11.813 8.25e-13 ***
## mpg          -8.83       1.31  -6.742 1.79e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 43.95 on 30 degrees of freedom
## Multiple R-squared:  0.6024, Adjusted R-squared:  0.5892
## F-statistic: 45.46 on 1 and 30 DF,  p-value: 1.788e-07
```

By looking at the numerous r-squared values, Chris is right mpg had a high r square value of 60% comp

2b)

```
summary(model2<-lm(hp~cyl+mpg,data = mtcars))
```

```
##
## Call:
## lm(formula = hp ~ cyl + mpg, data = mtcars)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -53.72 -22.18 -10.13  14.47 130.73
##
## Coefficients:
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept)  54.067      86.093   0.628  0.53492
## cyl          23.979       7.346   3.264  0.00281 **
## mpg          -2.775       2.177  -1.275  0.21253
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 38.22 on 29 degrees of freedom
## Multiple R-squared:  0.7093, Adjusted R-squared:  0.6892
## F-statistic: 35.37 on 2 and 29 DF,  p-value: 1.663e-08
```

```
((model2$coefficients[2]*4)+model2$coefficients[1])+(model2$coefficients[3]*22)
```

```
##      cyl
## 88.93618
```

```
predict(model2,data.frame(cyl=4,mpg=22),interval = "prediction",level=0.85)
```

```
##      fit      lwr      upr
## 1 88.93618 28.53849 149.3339
```

3a)

```
library(mlbench)
```

```
## Warning: package 'mlbench' was built under R version 4.2.2
```

```
data(BostonHousing)
hos<-lm(medv~crim+zn+prratio+chas,data=BostonHousing)
summary(hos)
```

```
##
## Call:
## lm(formula = medv ~ crim + zn + prratio + chas, data = BostonHousing)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -18.282  -4.505  -0.986   2.650  32.656
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  49.91868    3.23497   15.431 < 2e-16 ***
## crim        -0.26018    0.04015   -6.480 2.20e-10 ***
## zn           0.07073    0.01548    4.570 6.14e-06 ***
## prratio     -1.49367    0.17144   -8.712 < 2e-16 ***
## chas1        4.58393    1.31108    3.496 0.000514 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 7.388 on 501 degrees of freedom
## Multiple R-squared:  0.3599, Adjusted R-squared:  0.3547
## F-statistic: 70.41 on 4 and 501 DF,  p-value: < 2.2e-16
```

```
## Due to the model's extremely low R square value of 36%, it is not particularly accurate.
```

3bi)

```
summary(hos1<-lm(medv~chas,data = BostonHousing))
```

```
##
## Call:
## lm(formula = medv ~ chas, data = BostonHousing)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -17.094  -5.894  -1.417   2.856  27.906
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  22.0938     0.4176  52.902 < 2e-16 ***
## chas1         6.3462     1.5880   3.996 7.39e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 9.064 on 504 degrees of freedom
## Multiple R-squared:  0.03072,    Adjusted R-squared:  0.02879
## F-statistic: 15.97 on 1 and 504 DF,  p-value: 7.391e-05
```

```
hos1$coefficients
```

```
## (Intercept)      chas1
##  22.093843    6.346157
```

```
(hos1$coefficients[2]*0)+hos1$coefficients[1]
```

```
##      chas1
## 22.09384
```

```
(hos1$coefficients[2]*1)+hos1$coefficients[1]
```

```
## chas1
## 28.44
```

The home with chas of 1 is more expensive than the house without chas of 0 with a value of 4.3 util.

3bii)

```
summary(hos2<-lm(medv~ptratio,data = BostonHousing))
```

```
##
## Call:
## lm(formula = medv ~ ptratio, data = BostonHousing)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -18.8342  -4.8262  -0.6426   3.1571  31.2303
##
```

```
## Coefficients:
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept)  62.345      3.029   20.58  <2e-16 ***
## ptratio      -2.157      0.163  -13.23  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 7.931 on 504 degrees of freedom
## Multiple R-squared:  0.2578, Adjusted R-squared:  0.2564
## F-statistic: 175.1 on 1 and 504 DF,  p-value: < 2.2e-16
```

```
(hos2$coefficients[2]*15)+hos2$coefficients[1]
```

```
## ptratio
## 29.987
```

```
(hos2$coefficients[2]*18)+hos2$coefficients[1]
```

```
## ptratio
## 23.51547
```

*##Using the correlation coefficients, it can be seen that the house price declines as the ptratio rises
The cost of the home with the ptratio of 15 is more expensive than the cost of the house with the pt.*

3c)

```
summary(hos)
```

```
##
## Call:
## lm(formula = medv ~ crim + zn + ptratio + chas, data = BostonHousing)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -18.282  -4.505  -0.986   2.650  32.656
##
## Coefficients:
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept)  49.91868    3.23497   15.431  < 2e-16 ***
## crim        -0.26018    0.04015   -6.480 2.20e-10 ***
## zn           0.07073    0.01548    4.570 6.14e-06 ***
## ptratio     -1.49367    0.17144   -8.712  < 2e-16 ***
## chas1        4.58393    1.31108    3.496 0.000514 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 7.388 on 501 degrees of freedom
## Multiple R-squared:  0.3599, Adjusted R-squared:  0.3547
## F-statistic: 70.41 on 4 and 501 DF,  p-value: < 2.2e-16
```


If your p-value is low (0.05), you can reject the null hypothesis. The model summary shows that each

3d)

```
anova(hos)
```

```
## Analysis of Variance Table
```

```
##
```

```
## Response: medv
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
## crim	1	6440.8	6440.8	118.007	< 2.2e-16 ***
## zn	1	3554.3	3554.3	65.122	5.253e-15 ***
## ptratio	1	4709.5	4709.5	86.287	< 2.2e-16 ***
## chas	1	667.2	667.2	12.224	0.0005137 ***
## Residuals	501	27344.5	54.6		

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
## by comparing the p values
```

```
## 1) crim
```

```
## 2) ptratio
```

```
## 3) zn
```

```
## 4) chas
```